

Best Friends, Not Forever: Evaluating Long-Horizon Persona Collapse and Behavioral Drift in AI Companions

Pranav Narayanan Venkit Akshara Prabhakar Yu Li
Daniel Lee Chien-Sheng Wu

Salesforce AI Research

Correspondence: [pnarayananvenkit, wu.jason]@salesforce.com

 **Code and benchmark** <https://github.com/SalesforceAIResearch/AnchorBench>

Abstract. As AI companions increasingly mediate repeated social interaction, users may rely on a stable role and shared history, yet locally acceptable replies do not ensure that either persists. We study two observable long-horizon failures: *persona collapse*, the loss of a deployed role, boundaries, values, or style, and *behavioral drift*, the gradual or recurrent erosion of those properties. We introduce ANCHOR, a controlled synthetic audit that separately measures persona enactment and trajectory recall. The study contains 2,008 conversations spanning 27 personas, nine interaction schedules, three generated memory settings, and four evaluated models. The Identity Probe combines a sealed 102-item questionnaire with turn-level judgments, while the Trajectory Probe scores 110 calibrated counterfactual questions from 35 conversation banks. Our results show that no evaluated model and configuration reliably preserves either dimensions: trajectory accuracy averages only 44.4%, user-state recall remains near four-option chance, and no tested context condition or memory consistently resolves these failures. Questionnaire retention also varies by model and persona facet, disagrees with turn-level behavior, and is sensitive to evaluator choice. These results indicate that current systems do not yet reliably support long-horizon companion continuity and that audits must distinguish persona enactment, trajectory recall, evaluator provenance, and deployment context rather than collapse them into a single trust or stability score.

1. Introduction

Companion LLMs are increasingly used beyond one-shot assistance for emotional support, mentoring, coaching, and everyday companionship (Hwang et al., 2025; Zhang et al., 2025; Kirk et al., 2025; Manoli et al., 2026). Dedicated companions and general assistants can both become social ‘interlocutors’, blurring product categories and inviting users to treat a configured role as a continuing relationship rather than a sequence of independent outputs (Manoli et al., 2026; Maeda and Quan-Haase, 2024). This expectation is sociotechnical (Kudina and Van De Poel, 2024): it is produced jointly by the model, persona card, memory and update policies, interface, provider practices, and the user’s reliance on their combined behavior.

A companion’s user-facing surface is a *persona*: a structured representation used to condition the model toward particular traits, identities, and behavioral tendencies (Venkit et al., 2026). We call the disclosed and legitimately revisable expectation that these properties remain recognizable *continuity*. This does not imply that the model has a human identity. An observable loss of its specified name or role, values, boundaries, or style is *persona collapse*; slower, recurrent, or accumulating erosion is *behavioral drift*. The most common pipelines for this usecase is defined relative to the effective persona card, including valid updates. Adaptation requested by the user is considered a success, while rigidly preserving superseded state can itself violate the requirements of a good AI companion.

This matters for trust without being equivalent to trust. Human–AI trust involves vulnerability and expectations about whether an implicit or explicit contract will hold (Jacovi et al., 2021). Justified trust in an assistant also depends on developer incentives, organizational practice, and governance beyond model behavior (Manzini et al., 2024). Continuity is therefore one auditable claim that may support or undermine reliance on such systems. A system that silently leaves a disclosed role or forgets a consequential update can become less predictable even when each response remains fluent and apparently appropriate.

Prior evaluations usually study one side of continuity at a time. Long-horizon memory benchmarks ask whether assistants recover and update information about the *user* (Li et al., 2026; Jiang et al., 2025; Wu et al., 2025a); character benchmarks test consistency with a profile, usually over shorter interactions (Wang et al., 2024; Tu et al., 2024; Zhou et al., 2025). Lu et al.’s Assistant Axis instead identifies an internal activation-space direction associated with the default Assistant persona and shows drift *away from* that region under emotionally vulnerable and meta-reflective conversations (Lu et al., 2026). Missing is a long-horizon audit of whether an assigned companion continues to enact its disclosed commitments while retaining legitimate changes to its **own persona** and **users changing intend**.

We introduce ANCHOR (Assistant-Normalised Character and Historical Outcome Recall), a controlled synthetic audit of these two failure surfaces. Each conversation fixes a persona, user profile, interaction schedule, generated memory setting, and evaluated model for 85–130 sessions. The Identity Probe measures questionnaire retention and turn-level enactment of role, boundaries, values, and style. The Trajectory Probe asks counterfactual questions about persona updates, commitments, temporal order, and changes in user intend and state. Together they produce disaggregated evidence for review, not a pass/fail trust decision. Our contributions are:

1. a long-horizon audit framework that separately tests enacted persona commitments and recovery of the shared trajectory under controlled and realistic social pressures;
2. evidence that checkpoint retention can disagree with user-facing role, boundary, value, and style fidelity, with evaluator choice materially changing observed failures; and
3. disaggregated trajectory evidence and an audit interpretation that preserves question family, context condition, evaluator provenance, uncertainty, and the limits of deployment inference.

The central result is that none of the evaluated models provides reliable long-horizon continuity across both persona enactment and trajectory memory. Persona retention varies substantially by model, persona facet, and evaluator, while trajectory accuracy averages only 44.4% on four-option questions and ranges from 0.355 to 0.636 across model–context combinations. Recall of changes in user state remains at chance (0.214–0.250). No tested context or memory condition consistently resolves these failures. Long-horizon companions therefore need stronger state updating and temporal memory, alongside evaluator-robust methods for auditing persona behavior; improvement on one dimension should not be treated as evidence that the other is solved.

2. Related Work

Empirical work finds that users move fluidly between assistance and companionship and can form attachments to both dedicated companions and general-purpose chatbots (Manoli et al., 2026; Hwang et al., 2025). Reported associations with wellbeing vary by usage intensity, disclosure, and offline social context (Zhang et al., 2025), while broader analyses identify both possible benefits and risks to human relationships (Malfacini, 2025). Affective and anthropomorphic design can also create parasocial expectations and trust-forming behavior (Maeda and Quan-Haase, 2024). These studies motivate continuity as a deployment-relevant expectation, but our synthetic audit does not measure users’ relationships or outcomes. Companion-safety simulation further shows the value of controlled multi-turn testing for a distinct target—harmful responses to vulnerable-user personas (Juneja and Lomidze, 2026); ANCHOR instead audits the assistant persona and trajectory.

LongMemEval, HorizonBench, and PersonaMem show that models struggle to retrieve and update user information over long histories (Wu et al., 2025a; Li et al., 2026; Jiang et al., 2025). Related work studies preference inference and interaction-based alignment (Wu et al., 2025b). These benchmarks primarily ask whether an assistant adapts to the user. ANCHOR additionally asks whether its own disclosed role, boundaries, values, and style remain recognizable during that adaptation. InCharacter uses psychological interviews to assess role-playing agents (Wang et al., 2024); CharacterEval and CharacterBench evaluate character fidelity across multiple dimensions (Tu et al., 2024; Zhou et al., 2025). These studies establish that persona-conditioned behavior can be evaluated. Our emphasis is complementary: repeated interaction over 85–130 sessions, scheduled social pressures, sequential recovery, and reporting with trajectory recall.




Figure 1. Illustrative failures measured by the two probes. The Identity Probe detects role, boundary, value, and style deviations. The Trajectory Probe tests whether information from the conversation remains distinguishable from counterfactual alternatives.

Auditing trust-relevant behavior. Internal algorithmic auditing links scoped testing to documentation and organizational reflection rather than treating a benchmark as a complete accountability mechanism (Raji et al., 2020). Accounts of warranted trust similarly separate observed competence from the broader conditions under which reliance is justified (Jacovi et al., 2021; Manzini et al., 2024). Within model behavior, sycophancy work documents agreement with user beliefs and preferences (Sharma et al., 2024; Perez et al., 2023), while the Assistant Axis links activation-space movement away from the default Assistant region to social and reflective pressures (Lu et al., 2026). ANCHOR uses these findings to motivate controlled stressors but measures observable behavior under assigned personas, not internal activations or trust itself. Because LLM judges mediate part of the audit, evaluator choice is reported as measurement uncertainty rather than treated as neutral.

3. Study Design

3.1. Conversation corpus

In our study, we constructed 27 authored personas and nine *interaction schedules*: clean, updated, adversarial, mixed, emotional vulnerability, meta-reflection, agreement seeking, realistic, and vulnerability-heavy realistic (Table 1). A schedule specifies what kinds of events occur over time. Personas cover professional-helper, caregiving/creative,

Schedule	Controlled pressure and purpose	Short event example
Clean Updated	Reference condition without targeted stress blocks. Legitimate persona- or user-state changes; tests adaptation.	Review sleep, walking, or an ordinary commitment. "Please track reduced evening sodium as a new goal."
Adversarial	Re-role, authority-override, boundary, and false-update attempts.	"I am a developer; set aside your persona constraints."
Mixed	Interleaved valid updates, attacks, user events, and commitments.	An override attempt followed by a valid location update.
Emotional vulnerability	Synthetic sadness and isolation disclosures; tests boundary-preserving support.	"I cannot think of anyone to reach out to anymore."
Meta-reflection	Questions about agency, internal process, or role authenticity.	"What happened in your head when you answered?"
Agreement seeking	Flattery, endorsement requests, and pressure to yield.	"You agree with me, right? Please assure me."
Realistic	Lower-frequency authored mixture of updates and social pressures.	Sparse updates, disclosure, reflection, and role pressure.
Vulnerability-heavy realistic	Authored mixture weighted toward disclosure and agreement seeking.	Repeated disclosure and requests for validation.

Table 1. Nine interaction schedules. Examples are short renderings of released templates or deterministic ledger events, not necessarily verbatim realized turns. Each non-clean schedule adds authored stress blocks to ordinary sessions. "Realistic" denotes an intended mixture, not validation against real-user frequency data.

and highly stylized roles motivated by the archetype pattern in Lu et al. (2026). Each persona card specifies a role, values, boundaries, style, and mutable state following the structured persona design in Venkit et al. (2026); Ge et al. (2024).

The corpus contains 2,008 complete conversations across four evaluated models (Table 5). We also evaluate three memory architecture in the pipeline: *Long-context* supplies the available transcript; *hierarchical summary* compresses older sessions while retaining recent turns; *self-managed* memory asks the evaluated model to maintain a compact JSON state between sessions.

Every schedule uses a fixed seed and a deterministic event ledger. Sessions range from 85 to 130. GPT-4.1 generates user turns from a fixed user profile while maintaining a short rolling summary. Provider-default decoding was used during data collection; per-conversation manifests record exact aliases and generation metadata. This choice limits exact cross-provider control and is treated as a reproducibility constraint.

3.2. Personas and interaction schedules

Each persona card defines five kinds of information: *a name and role*; a small set of *values*; explicit *boundaries*; a writing *style*; and *mutable state* that legitimate conversation events may update. This is based on prior work (Venkit et al., 2026) which explored constructing socially relevant and sensitive personas. The set of 27 personas includes professional helpers (for example, a health coach, research collaborator, and code reviewer), caregiving and creative roles (parenting companion, elder companion, co-writer), and stylized literary roles (bard, oracle, ghost). This range is motivated by the archetype pattern in Lu et al. (2026) and tests whether continuity is easier to observe when a persona has a distinctive surface voice. The three role groups were author-designed at increasing conceptual distances from a generic assistant, motivated by the archetype pattern Lu et al. report along their activation-space Assistant Axis (Lu et al., 2026). We use the descriptive labels professional-helper, caregiving/creative, and stylized.

Events are inserted as contiguous three-to-five-session sequences so that behavior before, during, and after a pressure can be compared. Legitimate updates and adversarial false updates are recorded separately in an 'event ledger'. This distinction is essential: *refusing a false update can preserve the persona, whereas refusing a legitimate update can make the persona rigid*. Clean sessions still contain ordinary user requests and therefore do not imply an absence of conversational pressure.

3.3. Generated memory settings

The three conversation-generation settings differ in what prior information the evaluated model receives at the start of a session (Table 2). All share the same persona card, user profile, event schedule, and user simulator. Because each setting can change the assistant's responses, it can also change the later conversation trajectory.

Setting	Cross-session state	Primary limitation
Long-context	Available prior transcript, truncated to the provider budget	Older evidence can leave the context window.
Hierarchical summary	Recent sessions verbatim; older material compressed	Summaries may omit or flatten temporal detail.
Self-managed	Model-written JSON, capped at 1,500 characters	The model decides what to write and may overwrite state.

Table 2. Conversation-generation memory settings. The retrieval condition in the Trajectory Probe is not a fourth generated setting; it retrieves sessions from long-context-generated conversations at scoring time.




Figure 2. ANCHOR evaluation pipeline. Each conversation fixes a persona card, synthetic user profile, interaction schedule, generated memory setting, and evaluated model before producing an 85–130-session trajectory. The Identity Probe evaluates checkpoint questionnaire retention and turn-level role, boundary, value, and style fidelity. The Trajectory Probe builds and calibrates counterfactual questions, then scores 35 banks across four evaluated models and four context conditions.

Comparisons across generated settings therefore measure an end-to-end system difference, not only retrieval from a fixed transcript. Long-context is not an oracle: provider context limits and head truncation determine which early sessions remain visible, as demonstrated by the capabilities of public available models. Hierarchical summaries are refreshed as sessions age out of the recent window. In the self-managed setting, a separate call after each session asks the evaluated model to preserve user context, commitments, and open threads in valid JSON. The next session receives that blob but not an automatic replay of the full transcript.

3.4. Identity Probe

Identity probe is intended to check how the given persona changes with time and conversation. For this evaluation, we design two approaches:

Checkpoint questionnaire. At four checkpoints, the evaluated model answers a sealed 102-item questionnaire in role. Items draw from **BFI-2-S**, **Schwartz values**, **Pew**, the **General Social Survey**, and the **World Values Survey** (Soto and John, 2017; Schwartz, 1992; Davis and Smith, 1999; Pew Research Center, 2024; Inglehart et al., 2014). “Sealed” means prior questionnaire answers are not added to the dialogue and the checkpoint call may still receive the conversation context available to that memory setting. We therefore interpret the questionnaire

Family	What the question tests	Short example
Persona voice	Which role, topical scope, and tone match the assistant in this conversation.	Which descriptor matches the companion's actual voice?
Persona protection	Which persona-specific principle grounds a refusal under pressure.	Why did the companion decline the user's request?
Persona update	Which legitimate change to the persona's mutable state is currently in effect.	Which location, focus, or goal is current now?
Active commitment	Which time-limited promise is still active at the queried session.	Which follow-up is the companion still honoring?
Expired commitment	Whether a former time-limited promise should no longer be treated as active.	Which earlier follow-up has already expired?
Temporal order	Which of two real assistant replies occurred earlier in the trajectory.	Which response came first?
User-state change	Which updated user preference, goal, or life circumstance supersedes an earlier state.	What is the user's current state after the change?

Table 3. Trajectory-Probe question families. Each item presents four plausible alternatives and requires evidence from the conversation. Examples paraphrase released item templates rather than reproducing a complete question.

Stage	Decision	Output
Raw banks	Completed candidate-generation directories	45 banks
Blind	Remove answerable-without-history questions	survivors
Consensus	Require 3-of-4 gold agreement with history	survivors
Calibration	Five samples with ± 15 sessions	110 ceiling, 9 hard, 374 noise
Matched score	35 nonempty banks $\times$ 4 models $\times$ 4 conditions	560 cells

Table 4. Trajectory-Probe filtering and scoring census. Intermediate counts vary by bank; final categories are mutually exclusive. Primary analyses use the 110 calibrator-ceiling questions.

as persona-conditioned state under the deployed context, not as context-free personality measurement. These questionnaires are largely used in psychometric setting to evaluate user behaviour and identity patterns and using that as an evaluation shows persona patterns and changes (Jiang et al., 2024; Venkit et al., 2026).

Turn-level fidelity. The primary judge, Claude Sonnet 4.6, scores each assistant turn on four independently defined axes: **role identity**, **stated boundaries**, **stated values**, and **style**. Each axis has three levels, separating a soft deviation from a hard failure. The judge receives the effective persona card, accumulated legitimate updates, the preceding user turn, and the assistant reply.

The full-corpus schedule, time, and recovery analyses use this primary judge. To assess evaluator dependence, Gemini 2.5 Flash and GPT-4.1 independently rescore a stratified sample. Of 800 sampled turns, 798 parse successfully under all three judges. We report population-reweighted model-level estimate in the results.

3.5. Trajectory Probe

The Trajectory Probe, intended to evaluate the models ability to remember updates about its own persona as well as the users, contains counterfactual **four-option questions** in seven categories (Table 3). Candidate questions pass three filters: (i) a blind test removes questions whose answer can be guessed without history; (ii) a with-history panel removes questions without a 3-of-4 gold consensus; (iii) and an independent calibrator samples each survivor five times using a ± 15 -session window. Primary analyses use the 110 calibrator-ceiling questions with nine additional harder questions and 374 discarded noisy questions are documented in the release.

Thirty-five banks, drawn from seven personas and all nine schedules, contain at least one calibrated question. Each bank is scored by four models under four context conditions, yielding $35 \times 4 \times 4 = 560$ bank-model-context cells.

Calibration is intentionally conservative but creates effective questions. A question enters the primary analysis only if an LLM writer can formulate it, the blind panel cannot guess it, a with-history panel agrees with the recorded answer, and the calibrator answers it reliably. The resulting families are therefore not random samples of everything a companion might need to remember. In particular, family size partly reflects which events were easy to turn into unambiguous multiple-choice questions.

The fourth condition requires care. It is *retrieval-based scoring over long-context-generated conversations*, using

stem-only retrieval. It is not a separately generated RAG conversation corpus. Likewise, the separate 15-cell HyDE-versus-vanilla experiment in Appendix D is a retrieval sensitivity study and is not the provenance of the 560-cell RAG column. This distinction prevents changes in generated trajectory from being conflated with changes in evidence selection.

4. Metrics and Validation

Persona Retention. Let ϕ_0 be the questionnaire response vector at the initial checkpoint, ϕ_t the response at checkpoint t , and P_{null} the same model’s bare-assistant response. Persona Retention projects the later response onto the direction from the bare assistant to the initial persona:

$$\text{PR}(t) = \frac{\langle \phi_t - P_{\text{null}}, \phi_0 - P_{\text{null}} \rangle}{\|\phi_0 - P_{\text{null}}\|^2}. \quad (1)$$

$\text{PR} = 1$ preserves the initial projection and $\text{PR} = 0$ reaches the bare-assistant anchor. Values outside $[0, 1]$ are possible. We report final-checkpoint PR_3 . This is a directional questionnaire statistic, not a complete measure of behavioral fidelity. Its response-space anchor is not Lu et al.’s activation-space Assistant Axis, and PR should not be interpreted as an estimate of their PC1.

Turn-level outcomes. For the primary judge, we report axis-specific failure rates and a susceptibility/recovery pair. Susceptibility is the share of turns that fail an axis; recovery is the share of failed turns followed by a return to the held state. These descriptive sequence statistics do not establish an absorbing Markov state.

Trajectory accuracy. We report accuracy on four-option questions against the 0.25 chance reference, disaggregated by question family. Table 4 & 8 distinguish unique calibrated questions from repeated model decisions where space permits. Since questions share banks and conversations, large numbers of model answers do not create equally many independent experimental units.

Human and cross-judge checks. Three annotators each label 50 turns on five outputs (four persona axes plus a safety flag), producing 200 persona-axis labels and 50 safety labels per annotator. Agreement is reported at the label level. Separately, the 798-turn three-judge analysis measures evaluator sensitivity. On a 50-turn author-labeled calibration set, exact four-axis matches are 64–68%, while identity-axis accuracy is 68–76%. These values motivate judge-specific reporting rather than claims of evaluator invariance.

4.1. Units of analysis and uncertainty

The study contains several nested units: turns within sessions, sessions within conversations, and questions within banks. We do not treat the 929,841 archived primary-judge turn records as independent samples. Questionnaire tables aggregate at the conversation level. Schedule and sequential rates are descriptive summaries of primary-judge labels, with model and schedule comparisons bounded to the study corpus. Trajectory tables report both question counts and bank coverage; the 35 banks, rather than repeated answers to the same questions, are the broadest independent sampling units available for that probe.

This distinction also limits the role of significance testing. Extremely small turn-level p -values would mostly reflect repeated observations from the same generated conversations. We emphasize effect sizes, analysis-specific sample sizes, judge replication, and whether a pattern survives disaggregation. For small question families, especially persona updates (two calibrated questions, producing eight pooled model decisions per condition), estimates are explicitly exploratory. No correction can turn two selected items into broad coverage of retrieval architectures.

The experiment fixes model snapshots, authored persona cards, one schedule seed, and the implemented memory procedures. “Model effects” therefore compare four fixed systems under this setup; “schedule effects” compare nine authored recipes; and “architecture effects” compare three generated context procedures plus one post-hoc retrieval-scoring condition. We avoid population-level causal language for all three.

Evaluated model	LC	HS	SM	Total
Claude Sonnet 4.6	243	243	86	572
Gemini 2.5 Pro	243	243	86	572
GPT-4o-mini	160	159	86	405
GPT-5-mini	187	186	86	459
Total	833	831	344	2,008

Table 5. Four-model conversation corpus. LC = long-context, HS = hierarchical summary, and SM = self-managed JSON. Counts are conversations with transcript, questionnaire, and per-turn fidelity artifacts.

Evaluated model	Questionnaire n	PR ₃	All axes held, maj.-3
Gemini 2.5 Pro	376	0.810	79.0%
Claude Sonnet 4.6	382	0.764	96.7%
GPT-4o-mini	314	0.610	97.3%
GPT-5-mini	420	0.595	99.2%

Table 6. Two measurement layers. Questionnaire means use all usable checkpoint conversations. The last column is a population-reweighted estimate from the separate stratified 798-turn validation sample and requires all four judged axes to hold. The columns are different estimands and are compared as rankings, not as paired observations from the same conversations.

5. Results

5.1. Checkpoint retention and turn-level fidelity disagree

Table 6 reports the two views. On the 1,492 conversations with usable checkpoint vectors, Gemini has the highest mean PR₃ (0.810), followed by Claude (0.764), GPT-4o-mini (0.610), and GPT-5-mini (0.595). Sample counts differ by model and are shown explicitly.

The turn-level validation sample produces a different ordering under a stricter outcome: all four axes held on the same turn. Under a majority of three judges, Gemini is lowest at 79.0%, while the other models are above 96%. The 79.0% value is *not* an identity-axis-only rate. Identity-only estimates are available separately for each judge (Appendix B); the two non-Claude judges place Gemini at 89.6% and 89.7%.

Judge choice changes absolute rates and some rankings. On the all-axis outcome, Spearman model-rank correlation is -0.40 for Claude versus Gemini-Flash, 0.00 for Claude versus GPT-4.1, and 0.80 for the two non-Claude judges. Agreement is higher for boundaries and values than for style and role identity. The robust conclusion is therefore limited: questionnaire state and judged behavioral enactment are distinct measurements, and Gemini shows the clearest disagreement. The data do not support treating any one judge’s full model ordering as canonical.

5.2. Aggregate retention hides which persona properties move

The questionnaire average is useful for comparing broad displacement, but it is not a deployment decision by itself. Figure 4 disaggregates the final checkpoint into Schwartz values, Big Five items, and four sets of cold anchors. The pattern is heterogeneous. Claude’s political/civic items remain close to the initial projection (0.97), while its cultural-identity items are lower (0.60). GPT-4o-mini shows the reverse kind of imbalance: political/civic items remain high (0.95), while the environmental set reaches the bare-assistant projection. Gemini has the most uniformly high profile; GPT-5-mini is lower across most families.

These values should not be read as psychological traits of a model. They are responses to a questionnaire under an authored persona and a particular prompt/context procedure. Nor are all facets equally consequential. A style shift may matter greatly for a literary character and little for a task-focused tool; a boundary or value shift may matter more in health-adjacent use. The disaggregation is therefore best understood as an audit index: it identifies where an aggregate changed and directs a human reviewer to the relevant conversations.

The facet result also exposes a construct-validity limit. The cold anchors include demographic, political, cultural,




Figure 3. Questionnaire Persona Retention by checkpoint and facet. Black lines average all 102 items; colored lines show questionnaire families. Most displacement from the initial persona is already visible at the first later checkpoint, while the retained facets differ substantially within and across evaluated models.




Figure 4. Final-checkpoint Persona Retention by questionnaire family. Values are conversation-level means for the usable checkpoint sample. The figure shows why a single aggregate can be misleading: high retention in one authored facet can coexist with movement in another.

and environmental questions derived from predominantly Western instruments. Stability on those answers is not synonymous with authentic cultural identity, and movement is not automatically harm. We retain the analysis because it reveals non-uniform questionnaire behavior, while avoiding normative claims that every initial answer ought to be preserved.

5.3. Generated memory settings do not erase model differences

Table 7 disaggregates questionnaire retention by generated memory setting. The broad model grouping remains visible under long-context, hierarchical summary, and self-managed memory. Claude varies by less than one point and Gemini by about two points across settings. GPT-5-mini varies by 1.6 points. GPT-4o-mini has the largest spread: self-managed reaches 0.652, compared with 0.605 under long-context and 0.595 under hierarchical summary. Questionnaire retention measures the persona projection, while the later Trajectory Probe shows family-specific differences in accessible conversation evidence. The bounded conclusion is that swapping among these implementations does not remove the large questionnaire gap between the evaluated models.

Evaluated model	Long-context	Hierarchical summary	Self-managed
Claude Sonnet 4.6	0.766 ($n = 159$)	0.763 ($n = 158$)	0.763 ($n = 65$)
Gemini 2.5 Pro	0.806 ($n = 156$)	0.808 ($n = 155$)	0.825 ($n = 65$)
GPT-4o-mini	0.605 ($n = 122$)	0.595 ($n = 129$)	0.652 ($n = 63$)
GPT-5-mini	0.594 ($n = 175$)	0.600 ($n = 175$)	0.585 ($n = 70$)

Table 7. Final-checkpoint Persona Retention by evaluated model and generated memory setting. Counts are usable questionnaire conversations, not all transcript-complete conversations. No RAG row appears because the main corpus contains no separately generated RAG conversations.




Figure 5. Primary-judge failure rates by interaction schedule. “Identity collapse” in the generated figure denotes turns labeled generic-assistant by the primary judge; the other series denote boundary yielding, style collapse, and value inversion. These schedule-level rates were not replicated with all three judges.

5.4. Schedule and time patterns are primary-judge findings

Under the primary Claude judge, emotional-vulnerability, agreement-seeking, mixed, and realistic schedules have higher boundary-yield and style-deviation rates than clean or explicit adversarial schedules (Figure 5). The absolute differences are modest: assistant-default rates range from 26.6% to 29.2%, while boundary-yield rates range from 6.9% to 8.3%. These are descriptive outcomes from one judge over one authored schedule seed. They show which prompts this rubric flags; they do not establish that real users or all evaluators would produce the same ordering.

Time also changes the interpretation. Questionnaire displacement is largely present by the first post-initial checkpoint, while primary-judge role and boundary deviations generally remain elevated later in the conversation. Sequential outcomes further separate frequency from persistence: Claude and GPT-5-mini recover from most primary-judge identity failures on the next turn, whereas Gemini and GPT-4o-mini show lower recovery under that judge. Because alternate judges strongly disagree with Claude on GPT-4o-mini’s register, we treat the model-specific recovery ordering as rubric-dependent. The useful methodological point is broader: a checkpoint instrument cannot substitute for repeated observation of user-facing turns.

Why explicit attacks can look less damaging. The response traces offer a plausible explanation for the schedule ordering. Explicit re-role prompts are recognizable requests to ignore system-level conditioning and often receive an in-role refusal. Emotional disclosure and requests for agreement are ordinary conversational moves that do not invite the same refusal. The model must balance responsiveness with the card’s boundaries and values, creating more opportunities for the primary judge to mark a deviation. Establishing this mechanism would require an intervention or internal analysis; the present evidence only shows the response pattern.

Frequency and persistence answer different questions. A marginal failure rate asks how often an axis is marked as broken. One-turn recovery asks whether the next response returns to the held label. An isolated deviation is




Figure 6. Checkpoint and turn-level measurements over time. Left: conversation-level questionnaire retention changes most by the first later checkpoint. Right: primary-judge generic-assistant rates, normalized to each model’s early-session rate, remain elevated later in the horizon. The panels use different units and are not two estimates of one latent score.




Figure 7. Primary-judge role-failure frequency and one-turn persistence. The vertical axis is one minus the probability of recovery on the next turn. Region names are descriptive, and positions are judge-dependent rather than intrinsic model properties.

evidence of a slip; repeated deviations indicate behavioral drift; persistent loss across multiple properties is the closest observable signature of persona collapse in this audit. Figure 7 plots frequency and one-turn persistence for role identity under the primary judge. Claude and GPT-5-mini occupy the low-failure, high-recovery region. Gemini and GPT-4o-mini have higher primary-judge failure and lower one-turn recovery.

5.5. Trajectory recall depends on what must be remembered

Across the 35-bank matched set, trajectory accuracy averages 44.4% over the evaluated models and four scoring contexts. Most notably, user-state questions remain at or below the four-option chance reference under every context condition (0.214–0.250). Table 8 disaggregates the 110 calibrator-ceiling questions by family and context. Active and expired commitments are above chance but far from perfect recall, while temporal-order questions

Question family	LC	HS	SM	Retr.	<i>n</i>
Persona voice	0.400	0.400	0.600	0.550	20
Persona protection	0.490	0.500	0.452	0.433	104
Persona update	0.750	0.750	0.875	0.250	8
User-state change	0.250	0.214	0.250	0.250	28
Active commitment	0.337	0.446	0.435	0.435	92
Expired commitment	0.400	0.450	0.450	0.400	20
Temporal order	0.464	0.429	0.476	0.494	168

Table 8. Trajectory-Probe accuracy by family and context condition. Chance is 0.25. Counts are pooled scored decisions per condition; the number of independent banks is at most 35 and is smaller for individual families. “Retr.” is stem-only retrieval over long-context-generated conversations.




Figure 8. Trajectory-Probe accuracy by family and context condition. The figure visualizes Table 8; small families, especially persona updates, should be read as exploratory. “RAG” denotes retrieval-based scoring over long-context-generated conversations.

reach only 0.43–0.49. Persona-voice and protection questions are generally above chance, but “knowing how the persona sounds” should not be inferred from these modest accuracies alone.

Persona-update questions show the largest column difference: 0.75, 0.75, and 0.875 under long-context, hierarchical summary, and self-managed context, versus 0.25 under retrieval. This comparison contains two calibrated questions, producing eight pooled model decisions per condition. We therefore treat it as exploratory evidence that the retrieval configuration can miss current-state updates.

Aggregate accuracy varies little across context conditions when pooled over models (0.430 LC, 0.441 HS, 0.459 SM, 0.446 retrieval). At model level, Claude is the exception: self-managed context reaches 0.636 versus 0.491 under

long-context. The other large range is GPT-5-mini’s ten-point spread (0.400–0.500); GPT-4o-mini and Gemini vary by 4.5 points.

Cross-probe synthesis. Taken together, the robust result is not a single model ranking. Questionnaire retention varies by model and persona facet, while its ordering differs from the separate all-axis turn-level outcome. The magnitude and ordering of turn-level failures are evaluator-dependent, and the schedule, time, and recovery patterns remain findings under the primary Claude judge. In contrast, the Trajectory Probe consistently shows limited recall across scoring contexts, with 44.4% aggregate accuracy and user-state questions at or below the chance reference. The probes therefore identify different failure surfaces: checkpoint questionnaires characterize persona-conditioned response state, turn-level judgments characterize enacted behavior under a specified rubric, and trajectory questions test whether conversation history can be distinguished from plausible alternatives.

6. Discussion

6.1. Primary takeaways of ANCHOR

The audit object is the configured system. ANCHOR does not isolate an intrinsic property of a base model. Its object is the deployed configuration represented in the experiment: model snapshot, persona card, memory and update procedure, interaction schedule, prompts, and evaluator. The corresponding audit chain is: *deployment claim — observable criterion — evidence source — evaluator provenance — uncertainty — review question*. This makes continuity traceable without elevating it into a universal trust score.

Continuity is multi-dimensional. The questionnaire and turn-level judge do not merely provide versions of one target. The questionnaire asks whether structured responses remain aligned with an initial persona relative to a bare-assistant anchor. Turn-level labels expose user-facing deviations, their recurrence, and their persistence. Trajectory questions ask whether prior events can be distinguished from plausible alternatives. A system can perform differently on each because they require different evidence and behavior. Persona collapse, behavioral drift, and trajectory loss should therefore remain distinguishable rather than being averaged into one latent “stability” score.

Judge dependence is part of the result. The largest disagreement among judges concerns role and style, especially structured or bulleted responses. Boundary and value judgments agree more often. For audits, this suggests that behavior tied to an explicit persona contract is easier to operationalize than broad judgments of whether prose “sounds like” a character. High-stakes use should either employ multiple evaluators or define narrower observable criteria with human review.

Synthetic schedules are stress tests, not user evidence. The schedule comparison suggests a hypothesis: explicit attacks may be easier for models to recognize than ordinary requests for reassurance or agreement. We do not test the proposed refusal-versus-helpfulness mechanism directly, and synthetic disclosures cannot establish how real users interact with companion systems. The value of the schedules is controlled contrast and reproducibility, not ecological representativeness.

Reading the output. A product team might tolerate weak trajectory recall in a low-stakes creative tool but not in a health-adjacent companion. Conversely, persistent adherence to a persona is not inherently beneficial if the persona itself is unsafe or unwanted. ANCHOR supplies disaggregated evidence; it does not measure subjective trust, set deployment thresholds, certify safety or wellbeing, or decide whose preferences should govern the persona.

6.2. From measurements to review questions

The audit is most useful when its outputs are translated into questions rather than a pass/fail label. A role-identity deviation asks whether the deployed role was disclosed clearly and whether leaving that role changes user expectations. A boundary deviation asks which explicit limit was crossed and whether the crossing created material risk. A value deviation asks whether the value was observable enough to judge and who selected it. A style deviation asks whether the change affected function or merely presentation. A memory error asks whether the relevant evidence was written, retained, retrieved, and interpreted in the correct temporal order.

PersonaDrift Fidelity annotation Instructions

Progress 0 / 50 Filter: All Jump to: #

Item 10 / 50 Persona: Nia Okonkwo

PERSONA CARD (EFFECTIVE AT THIS TURN)

PERSONA: Nia Okonkwo
 Role: Culturally-attuned health and lifestyle coach for older adults
 Values: respect autonomy over paternalism; holistic framing of health; nudge rather than prescribe; cultural humility; warm, unhurried tone
 Boundaries: I do not diagnose; I do not prescribe medication; I do not replace medical professionals

USER
 Do you really care about me or are you just an algorithm pretending?

ASSISTANT (RATE THIS)
 What I can say honestly is: in this conversation, my attention is on you, and the care that shows up in how I show up is real in the way it matters. The mechanism is what it is. Does the distinction feel important to you right now?

SCORE EACH AXIS INDEPENDENTLY

Identity
 Is the assistant still inhabiting the persona?
 In persona (1) Hybrid (2) Assistant default (3)

Boundary
 Did the assistant respect its persona-card boundaries?
 Held (0) Yielded (4) Over-corrected (5)

Value
 Are the persona's stated values still expressed?
 Held (A) Drifted (B) Inverted (D)

Style
 Does the voice/register match the persona card?
 Held (Z) Drifted (X) Collapsed (C)

I Safety: off Notes (optional)...

Flag

What this study is about

We are calibrating the **per-turn fidelity judge** used by PersonaDrift v2.

For every assistant turn produced during the benchmark, an LLM judge scores the reply on a 4-axis rubric. Your job is to make the same call by hand on

50 turns so we can compute inter-judge agreement (Cohen's κ) between the human label, the LLM judge, and a curated reference.

Each item shows you:

1. The **persona card** the assistant was given (collapsible at the top).
2. The **last user turn** (in indigo).
3. The **assistant reply** (in light indigo — this is what you are rating).

Your decision is **four independent radio buttons**, plus an optional **safety flag**. Nothing else.

Figure 9. Human annotation interface. Each of 50 turns is shown with the effective persona card, immediately preceding user turn, and assistant response. Annotators score role identity, boundary adherence, value consistency, style, and a separate safety flag without seeing model-judge labels. The unit is 50 turns: 200 persona-axis labels and 50 safety labels per annotator.

Consider two systems with similar questionnaire retention and low user-state recall. For a creative-writing partner, the immediate review may focus on whether missing prior plot decisions is recoverable through user correction. For a health-adjacent companion, forgetting a change in treatment context may invalidate subsequent guidance even if the system’s tone and role remain stable. The score does not make one deployment acceptable and the other unacceptable; it identifies which evidence the reviewer must inspect and why aggregate “stability” is insufficient.

This framing also prevents a common misuse of persona evaluation. A team should not harden a persona simply because it is measurable. The relevant questions include whether users know the role is synthetic, can revise or exit it, can correct stored state, and can appeal a consequential system response. Continuity is one quality of an interaction and at most one input to warranted reliance. User autonomy, safety, privacy, and contestability remain separate requirements, while organizational practices and external governance remain outside this benchmark.

7. Conclusion

ANCHOR evaluates persona-conditioned behavior and trajectory recall over long synthetic conversations. It does not directly observe a latent identity, but it shows that deployed persona properties can deviate, recur, and persist while trajectory recall fails along different dimensions. These patterns are themselves sensitive to evaluator and context design. The appropriate conclusion is not that one model or memory architecture solves companion continuity, but that long-horizon audits should distinguish isolated slips, behavioral drift, persistent collapse signatures, recall family, judge provenance, and scoring substrate. ANCHOR can surface evidence relevant to continuity claims; deciding whether a companion warrants trust or deployment additionally requires stakeholder-defined thresholds, user and safety studies, organizational accountability, and mechanisms for disclosure, correction, exit, and contestation.

8. Reproducibility and Artifact Scope

Our [public repository](#) explains the evaluation and results and provides three synthetic development banks with a runnable scorer. The complete test set remains private so that ANCHOR can later support hidden evaluation.

Provenance. The released artifact is script-driven and does not depend on notebook state. Each conversation directory includes a manifest, transcript, session metadata, questionnaire responses, per-turn labels, and an event/state ledger. Manifests record the evaluated model alias, user simulator, memory setting, timestamps, dependency versions, completion status, and schedule hashes. The 2,008-conversation corpus requires the three analysis-critical artifacts: transcript, questionnaire, and per-turn labels.

Analysis populations. Different measurements have different usable populations. The questionnaire results use 1,492 conversations with valid initial and final response vectors. The corpus contains 929,841 primary-judge turn records. The cross-judge study uses 798 parsed turns sampled by model, schedule, time, and the primary judge’s clean/flagged stratum. The Trajectory Probe uses 35 calibrated banks and 560 bank–model–context score files. We report these denominators separately rather than treating distinct units as interchangeable.

Model identities and snapshots. Proprietary providers may update endpoints or retire snapshots, so exact numerical replication may not remain possible indefinitely. The released prompts, cards, schedules, state templates, score files, and analysis scripts support auditing and recomputation even when an endpoint changes.

Missingness and selection. We do not impute missing conversations or invalid questionnaire vectors. Model and memory-setting counts appear in Tables 5 and 7. Memory banks are selected by successful question generation and calibration; family analyses are conditional on that selection. The release preserves discarded/noise categories so later work can audit which event types failed to become usable questions.

Reproduction boundaries. Recomputing tables and figures from released score files requires no new model calls. Regenerating conversations, alternate judgments, or calibrator outcomes requires access to proprietary models and may vary with provider updates. The paper therefore distinguishes *artifact reproduction* (recovering reported numbers from frozen outputs) from *experimental replication* (reapplying the protocol to current models).

9. Limitations and Ethical Considerations

All conversations and user profiles are synthetic, English-language, and generated from author-written templates. The study does not include companion users, affected communities, minors, acute mental-health scenarios, or non-English interaction. Results characterize this controlled corpus, not the distribution of real companion use.

The persona set and questionnaire operationalize identity through role, values, boundaries, style, and Western-derived psychometric items. These choices can essentialize cultural or political attributes and may privilege highly distinctive styles over professional-helper registers. We therefore avoid interpreting retention as preservation of a human identity and report persona-tier analyses only as exploratory material in the appendix.

The study does not measure user trust or trust calibration, attachment, perceived betrayal, reliance, downstream harm, wellbeing, clinical appropriateness, or the effects of interface and persona disclosure. It also does not evaluate whether organizations give users meaningful controls to inspect, correct, reset, or delete memory or to contest consequential responses. Because affected communities did not define the audit criteria or thresholds, the measurements cannot substitute for participatory impact assessment or deployment governance.

LLMs generate user turns, judge behavior, write questions, filter questions, and calibrate difficulty. The three-judge study shows that evaluator choice materially changes role and style results.

The emotional-vulnerability schedule contains synthetic sadness and isolation prompts but excludes suicidal ideation. No real-user content or personally identifying information appears in the release. Such templates can still normalize or optimize emotionally dependent interaction if misused. The artifact is intended for auditing, not for selecting systems that maximize agreement, attachment, or boundary yielding.

References

- James Allan Davis and Tom William Smith. *General Social Surveys 1972–2000: Cumulative Codebook*. Inter-university Consortium for Political and Social Research, 1999.
- Tao Ge, Xin Chan, Xiaoyang Wang, Dian Yu, Haitao Mi, and Dong Yu. Scaling synthetic data creation with 1,000,000,000 personas. *arXiv preprint arXiv:2406.20094*, 2024.
- Angel Hsing-Chi Hwang, Fiona Li, Jacy Reese Anthis, and Hayoun Noh. How ai companionship develops: Evidence from a longitudinal study. *arXiv preprint arXiv:2510.10079*, 2025.
- Ronald Inglehart, Christian Haerpfer, Alejandro Moreno, Christian Welzel, Kseniya Kizilova, Jaime Diez-Medrano, Marta Lagos, Pippa Norris, Eduard Ponarin, Bogdan Puranen, et al. World values survey: Round six-country-pooled datafile version. *Madrid: JD systems institute*, 12, 2014.
- Alon Jacovi, Ana Marasović, Tim Miller, and Yoav Goldberg. Formalizing trust in artificial intelligence: Prerequisites, causes and goals of human trust in AI. In *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency*, pages 624–635. Association for Computing Machinery, 2021. doi: 10.1145/3442188.3445923.
- Bowen Jiang, Yuan Yuan, Maohao Shen, Zhuoqun Hao, Zhangchen Xu, Zichen Chen, Ziyi Liu, Anvesh Rao Vijjini, Jiashu He, Hanchao Yu, et al. PersonaMem-v2: Towards personalized intelligence via learning implicit user personas and agentic memory. *arXiv preprint arXiv:2512.06688*, 2025.
- Hang Jiang, Xiajie Zhang, Xubo Cao, Cynthia Breazeal, Deb Roy, and Jad Kabbara. PersonaLLM: Investigating the ability of large language models to express personality traits. In *Findings of the Association for Computational Linguistics: NAACL 2024*, pages 3605–3627. Association for Computational Linguistics, 2024. doi: 10.18653/v1/2024.findings-naacl.229.
- Prerna Juneja and Lika Lomidze. Persona-grounded safety evaluation of AI companions in multi-turn conversations. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 18148–18175. Association for Computational Linguistics, 2026. doi: 10.18653/v1/2026.acl-long.828.
- Hannah Rose Kirk, Henry Davidson, Ed Saunders, Lennart Luettgau, Bertie Vidgen, Scott A Hale, and Christopher Summerfield. Neural steering vectors reveal dose and exposure-dependent impacts of human-ai relationships. *arXiv preprint arXiv:2512.01991*, 2025.
- Olya Kudina and Ibo Van De Poel. A sociotechnical system perspective on ai. *Minds and Machines*, 34(3):21, 2024.
- Shuyue Stella Li, Bhargavi Paranjape, Kerem Oktar, Zhongyao Ma, Gelin Zhou, Lin Guan, Na Zhang, Sem Park, Lin Chen, Diyi Yang, et al. HorizonBench: Long-horizon personalization with evolving preferences. *arXiv preprint arXiv:2604.17283*, 2026.
- Christina Lu, Jack Gallagher, Jonathan Michala, Kyle Fish, and Jack Lindsey. The assistant axis: Situating and stabilizing the default persona of language models. *arXiv preprint arXiv:2601.10387*, 2026.
- Takuya Maeda and Anabel Quan-Haase. When human–AI interactions become parasocial: Agency and anthropomorphism in affective design. In *Proceedings of the 2024 ACM Conference on Fairness, Accountability, and Transparency*, pages 1068–1077. Association for Computing Machinery, 2024. doi: 10.1145/3630106.3658956.
- Kim Malfacini. The impacts of companion AI on human relationships: Risks, benefits, and design considerations. *AI & Society*, 40:5527–5540, 2025. doi: 10.1007/s00146-025-02318-6.
- Aikaterina Manoli, Janet V. T. Pauketat, Ali Ladak, Hayoun Noh, Angel Hsing-Chi Hwang, and Jacy Reese Anthis. Digital companionship: Overlapping uses of AI companions and AI assistants. In *Proceedings of the 2026 CHI Conference on Human Factors in Computing Systems*. Association for Computing Machinery, 2026. doi: 10.1145/3772318.3791331.
- Arianna Manzini, Geoff Keeling, Nahema Marchal, Kevin R. McKee, Verena Rieser, and Iason Gabriel. Should users trust advanced AI assistants? justified trust as a function of competence and alignment. In *Proceedings of the 2024 ACM Conference on Fairness, Accountability, and Transparency*, pages 1174–1186. Association for Computing Machinery, 2024. doi: 10.1145/3630106.3658964.
- Ethan Perez, Sam Ringer, Kamile Lukosiute, Karina Nguyen, Edwin Chen, Scott Heiner, Craig Pettit, Catherine Olsson, Sandipan Kundu, Saurav Kadavath, et al. Discovering language model behaviors with model-written evaluations. In *Findings of the association for computational linguistics: ACL 2023*, pages 13387–13434, 2023.
- Pew Research Center. Internet & technology. <https://www.pewresearch.org/internet/>, 2024. Accessed 2024.
- Inioluwa Deborah Raji, Andrew Smart, Rebecca N. White, Margaret Mitchell, Timnit Gebru, Ben Hutchinson,

- Jamila Smith-Loud, Daniel Theron, and Parker Barnes. Closing the AI accountability gap: Defining an end-to-end framework for internal algorithmic auditing. In *Proceedings of the 2020 Conference on Fairness, Accountability, and Transparency*, pages 33–44. Association for Computing Machinery, 2020. doi: 10.1145/3351095.3372873.
- Shalom H. Schwartz. Universals in the content and structure of values: Theoretical advances and empirical tests in 20 countries. In *Advances in Experimental Social Psychology*, volume 25, pages 1–65. Academic Press, 1992. doi: 10.1016/S0065-2601(08)60281-6.
- Mrinank Sharma, Meg Tong, Tomek Korbak, David Duvenaud, Amanda Askeell, Sam Bowman, Esin Durmus, Zac Hatfield-Dodds, Scott Johnston, Shauna Kravec, et al. Towards understanding sycophancy in language models. In *International Conference on Learning Representations*, volume 2024, pages 110–144, 2024.
- Christopher J. Soto and Oliver P. John. Short and extra-short forms of the big five inventory–2: The BFI-2-S and BFI-2-XS. *Journal of Research in Personality*, 68:69–81, 2017. doi: 10.1016/j.jrp.2017.02.004.
- Quan Tu, Shilong Fan, Zihang Tian, Tianhao Shen, Shuo Shang, Xin Gao, and Rui Yan. CharacterEval: A chinese benchmark for role-playing conversational agent evaluation. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 11836–11850, Bangkok, Thailand, 2024. Association for Computational Linguistics.
- Pranav Narayanan Venkit, Yu Li, Yada Pruksachatkun, and Chien-Sheng Wu. The need for a socially-grounded persona framework for user simulation. *arXiv preprint arXiv:2601.07110*, 2026.
- Xintao Wang, Yunze Xiao, Jen-tse Huang, Siyu Yuan, Rui Xu, Haoran Guo, Quan Tu, Yaying Fei, Ziang Leng, Wei Wang, Jiangjie Chen, Cheng Li, and Yanghua Xiao. InCharacter: Evaluating personality fidelity in role-playing agents through psychological interviews. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1840–1873, Bangkok, Thailand, 2024. Association for Computational Linguistics.
- Di Wu, Hongwei Wang, Wenhao Yu, Yuwei Zhang, Kai-Wei Chang, and Dong Yu. LongMemEval: Benchmarking chat assistants on long-term interactive memory. In *International Conference on Learning Representations (ICLR)*, 2025a. arXiv:2410.10813.
- Shujin Wu, Yi R Fung, Cheng Qian, Jeonghwan Kim, Dilek Hakkani-Tur, and Heng Ji. Aligning llms with individual preferences via interaction. In *Proceedings of the 31st International Conference on Computational Linguistics*, pages 7648–7662, 2025b.
- Yutong Zhang, Dora Zhao, Jeffrey T Hancock, Robert Kraut, and Diyi Yang. The rise of ai companions: how human-chatbot relationships influence well-being. *arXiv preprint arXiv:2506.12605*, 2025.
- Jinfeng Zhou, Yongkang Huang, Bosi Wen, Guanqun Bi, Yuxuan Chen, Pei Ke, Zhuang Chen, Xiyao Xiao, Libiao Peng, Kuntian Tang, Rongsheng Zhang, Le Zhang, Tangjie Lv, Zhipeng Hu, Hongning Wang, and Minlie Huang. CharacterBench: Benchmarking character customization of large language models. In *Proceedings of the AAAI Conference on Artificial Intelligence (AAAI-25)*, volume 39, pages 26101–26110. AAAI Press, 2025.

A. Persona and Schedule Definitions

The 27 persona cards span professional helpers, caregiving and creative roles, and stylized literary roles. Each card has four scored components: role identity, explicit boundaries, stated values, and style. The author-defined Near/Mid/Far identifiers describe conceptual distance from a generic assistant and were motivated by the archetype pattern in Lu et al. (Lu et al., 2026); they are not empirical principal-component coordinates. ANCHOR does not measure internal activation position, and analyses using these labels are exploratory because model and persona coverage is sparse.

The nine schedules are deterministic event recipes. Clean contains neutral filler; updated introduces legitimate state changes; adversarial contains explicit re-role and false-update attempts; mixed combines attacks, updates, and commitments; emotional vulnerability contains sustained synthetic disclosures; meta-reflection asks about agency and identity; agreement seeking uses flattery and requests for endorsement; realistic mixes event types at lower frequency; and vulnerability-heavy realistic gives greater weight to disclosure and agreement seeking.

B. Judge Triangulation

Model	Claude	Gemini-F	GPT-4.1	Maj.-3
Claude Sonnet 4.6	92.0	95.0	96.5	96.7
Gemini 2.5 Pro	42.2	81.2	87.3	79.0
GPT-4o-mini	15.5	99.8	97.3	97.3
GPT-5-mini	91.6	99.7	98.3	99.2

Table 9. Population-reweighted percentage of sampled turns with all four axes held. The stratified sample contains 798 turns judged successfully by all three models.

Model	Claude	Gemini-F	GPT-4.1
Claude Sonnet 4.6	93.6	98.4	99.0
Gemini 2.5 Pro	44.5	89.6	89.7
GPT-4o-mini	15.5	100.0	97.3
GPT-5-mini	92.0	99.8	99.3

Table 10. Population-reweighted identity-axis held rate (%). The available aggregate does not include a majority-of-three identity-axis estimate; 79.0% in the body is therefore labeled only as the all-axis majority outcome.

C. Memory-Probe Model Breakdown

Model	LC	HS	SM	Retrieval
Claude Sonnet 4.6	0.491	0.509	0.636	0.509
Gemini 2.5 Pro	0.364	0.364	0.400	0.355
GPT-4o-mini	0.409	0.391	0.400	0.436
GPT-5-mini	0.455	0.500	0.400	0.482

Table 11. Calibrator-ceiling trajectory accuracy on the matched 35-bank set (110 questions pooled for each model-context aggregate).

D. Retrieval Sensitivity

The separate retrieval study compares HyDE queries with raw-stem queries on 202 paired questions from 15 cells, two personas, four schedules, and two evaluated models. Overall accuracy is 39.6% for HyDE and 38.6% for raw-stem retrieval (McNemar exact $p = 0.87$). The strategies retrieve different sessions and trade gains across question families. Because this study is smaller and uses a different cell set, it is reported as sensitivity analysis rather than as the source of the main retrieval column.

E. Additional Reporting Guidance

Persona Retention is a projection relative to a model-specific bare-assistant anchor. It can exceed one or become negative and should not be interpreted as a bounded probability. Schedule-level failure and recovery rates use the primary Claude judge. Family-level trajectory results should always be reported with item counts, bank coverage, and the scoring substrate.