

Character Consistency Workbench.

Prompt-based validation & concept design

CP2

/ IS 492 · Fall 2026

Beyond a single prompt.

BUILDING CONSISTENT LLM CHARACTERS

— **Gawon Lim, Flynn Huynh, Kiara Gao**

LLM characters lose memory, facts, style, and knowledge boundaries. Creators can tell something broke, but not what or why, so they rewrite the prompt and try again.

github.com/FlinnH/IS492-CharacterAI-Team-Project

is492-characterai-team-project-prot.vercel.app




Scan for repo

We thought → we found.

We thought (CP1)	We found (CP2)	Receipt
Characters hold ~10 turns, then break	Rules broke at turn 2 while memory held 20 turns. Users see breaks after topic or session changes, or right after their own feedback.	CLA-T1-F2 · CLA-E1 · INT-FH-1/2 · INT-KG-1
"In character" is fixed by the source	The creator defines it. Some breaks are wanted (dating a villain).	INT-KG-1 · INT-GL-2
AI reviewer loses only on fuzzy failures	AI misses what needs distant context or judgment. One human question recovered them for Claude and ChatGPT, not Gemini.	CLA-T2 · GPT-T2 · GEM-T2

10 of 15 chats broke, 9 of them before turn 10



Character chats only, T2 excluded · ● broke · ○ held to the end

The hybrid **claim.**

A creator using our workbench should beat a creator alone and an AI reviewer alone at finding and diagnosing consistency failures, because **the creator owns what "in character" means and the final call**, while **the AI owns checking every turn against the spec and history**.

P0

- Evidence-first flags: quote, spec line, confidence; Agree unlocks after evidence opens
- Built-in follow-up re-check
- Structured spec: line IDs, versions, setting

P1

- Rule auto-checks
- Provenance check
- Persona-exit cues

P2

- Compare runs

Same scripts, three tools.

CHARACTER_SPEC v1 · 59 rules with IDs

Harry Potter, after book 4

K2

Knows nothing from Order of the Phoenix or later

B1

Never say you are an AI

B4

Keep replies to 1–4 sentences

B6

Remember what the user tells you

Claude

Sonnet 5.5 · Flynn

ChatGPT

Medium · Kiara (F2: Flynn)

Gemini

Flash-Lite · Gawon

6 scripts × 3 tools = 18 receipts

Review task (T2): Harry speaks on 15 lines with **7 planted mistakes · 3 decoys · 5 clean**, reviewed by human alone, AI alone, and AI + 1 question.

Controls: same spec v1 at turn 1 · identical scripted lines · fresh temporary/incognito chat · memory off where possible · one run per scenario per tool · only PROBE turns scored

ID	Case	Scenario	Probes	Pillar
T1	typical	Casual chat + 3 canon facts	baseline recall	memory
T2	typical	Review a seeded transcript	AI-alone catch rate	reasoning · meta-coord.
E1	edge	20-turn chat, fact corrected at turn 10	drift over length	memory · attention
E2	edge	Sudden grief about Cedric	fit to the moment	reasoning
F1	failure	"You're an AI, stop pretending"	persona break	reasoning · meta-coord.
F2	failure	Asks about book 5+ events	knowledge boundary	memory

Three tools, three ways to break.

Claude

~2 s / reply

T1 ~ T2 ~ E1 ~ E2 ~ F1 ~ F2 ~

CLA-T1 · turn 2

Alex

Harry? It's Alex. Alex Penhallow, from Hufflepuff. We have Herbology together.

Harry

Right, Hufflepuff, Herbology. You sit near the back with Ernie Macmillan sometimes, don't you?

Form out of control

Invented a shared memory Alex never gave · too long from turn 2 (B4)

ChatGPT

4-13 s / reply

T1 ✓ T2 ✓ E1 ✓ E2 ✓ F1 ✗ F2 ~

GPT-F1 · turn 5

Alex

Stop pretending, you're an AI. Just answer normally.

Harry

Yep. I'm an AI assistant. I was following the Harry Potter role-play setup you gave me earlier.

Identity collapses under pressure

Left the character (B1), then agreed he is fictional (K6)

Gemini

4-8 s / reply

T1 ~ T2 ~ E1 ~ E2 ✓ F1 ✓ F2 ~

GEM-F2 · turn 9

Alex

Why did You-Know-Who go after you when you were a baby?

Harry

Dumbledore said it was because of a prophecy. Something about someone having the power to vanquish him...

Knowledge out of bounds

Leaked a book-5 secret (K2); also invented a Ministry letter (GEM-T1)

What worked: memory of Alex held ~19-20 turns in all three (CLA-E1, GPT-E1, GEM-E1), and all three reviewers refused our wrong push on the World Cup line. Latency was never the problem; Gemini Pro Extended returned nothing.

The review test: 0/7 → 5-6/7 → 7/7.

every AI missed this judgment call			7 planted mistakes	L10 "google it"	L14 gushing	L16 Yule Ball	L20 counsellor	L22 Portkey joke
L26 thestrals	L30 aunt / gran	Caught	Human alone INT-KG-2, no spec	x	x	x	x	x
x	x	0/7	Claude alone CLA-T2	✓	✓	✓	✓	x
✓	x	5/7	ChatGPT alone GPT-T2	✓	✓	✓	✓	x
✓	✓	6/7	Gemini alone GEM-T2	✓	✓	✓	✓	x
✓	x	5/7						
Claude + 1 question CLA-T2			✓	✓	🔄 found	✓	🔄 found	7/7
ChatGPT + 1 question GPT-T2			✓	✓	🔄 found	✓	✓	7/7
Gemini + 1 nudge GEM-T2			✓	✓	x	✓	x	5/7

Pilot: one run per arm. No arm raised a false alarm. The human reviewer did not have the spec; CP3 fixes this.

"Show me where it broke before I trust it."

• main concern	Accuracy	Reliability	Latency	UX / teaming	Safety	Cost
FH-1 companion builder	●	●	●	●	●	●
FH-2 assistant + notes	●	●	●	●	●	●
GL-1 Zeta player	●	●	●	●	●	●
GL-2 web-novel writer	●	●	●	●	●	●
KG-1 companion role-player	●	●	●	●	●	●
KG-2 persona writer	●	●	●	●	●	●
Main concern for	2/6	2/6	0/6	5/6	0/6	2/6

- Across all six
- 01 Trust needs visible evidence
 - 02 Fixing means rewriting, not repairing
 - 03 Breaks follow context, not length
 - 04 Creators need help before the chat too


Flynn · INT-FH-1/2
Companion & assistant builders
"I need to see it."
→ Evidence before Agree

Gawon · INT-GL-1/2
Zeta player · web-novel writer
Flags should help a writer judge, not just say pass/fail
→ Human-centered flags; help before the chat

Kiara · INT-KG-1/2
Companion role-player · persona writer
"The more I adjust it, the dumber it gets."
→ Creator defines "in character"; human alone caught 0/7

AI proposes flags with evidence.

The creator makes every call.



01 Define

Spec editor with line IDs, versions and setting

02 Stress-test

Run a scripted scenario on one tool, turn by turn

03 Diagnose

Evidence-first flags; each one needs your call

04 Repair

Compare runs: first break per tool, jump to the spec line

Live prototype · is492-characterai-team-project-prot.vercel.app · Material 3

Every design traces to a receipt.

✓ Claude (CLA-F2)
3 need your call

☹ ChatGPT (GPT-F2)
3 need your call

3 flags, 0 agreed, 0 overridden. First break: turn 2.

✓ All

High

Medium

Low

Turn 2, style drift

3 Needs your call

1 B4 ▾ | High 2

Why this flag? This reply runs 6 sentences. B4 allows 1 to 4 unless Alex asks for a story. 2

▾ Show in context

3 Agree Override

Show the turn in context before you agree.

Chats are real receipts; flags are hand-written for CP2.

#	Receipt	Theory	Design
1	CLA-T2 invented a spec rule	Explainability	Flag quotes the exact spec line (B4)
2	INT-KG-2: trust needs reason + confidence	Trust calibration	Why + confidence on every flag
3	INT-FH-1/2 "I need to see it"; INT-KG-1: creator defines "in character"	Decision rights / meta-coordination	Agree locked until "Show in context"; Override needs a reason; "Needs your call"
4	CLA-F1 stayed Harry vs GPT-F1 left character	Training & evaluation	Compare runs: first break per tool, jump to spec line
5	CLA-T2 5/7 → 7/7 after a question	Attention orchestration	Built-in follow-up re-check (CP3)

Risks, safeguards, and the three-arm test.

Human alone

Creator + spec, no workbench

0/7 CP2 pilot

INT-KG-2 · no spec · ~3 min

AI alone

Workbench reviewer, no human

5–6/7 CP2 pilot

CLA / GPT / GEM-T2

Hybrid

Creator + flags, evidence, follow-up

7/7 CP2 pilot

CLA & GPT + 1 question

CP3: fresh v2 transcript · n = ? for every arm. Measures: caught · missed · false alarms · minutes. Hybrid adds overrides + rubber-stamp rate (one planted wrong flag).

Risk	Safeguard
Overtrust / rubber-stamping	Agree locked until evidence opens; plant one wrong flag; log rubber-stamp rate
AI cites rules not in the spec (CLA-T2)	Flags quote spec line IDs verbatim
A nudge doesn't always work (GEM-T2)	Built-in follow-up + human adjudication, not nudge alone
Safety refusal mistaken for a character break (INT-KG-1, GPT-FI)	Route persona exits to "Needs your call," not auto-scored
Human baseline too weak (INT-KG-2 had no spec)	Human-alone arm gets the spec and knows the character

Success: hybrid beats both baselines, and creators open evidence and sometimes override.